DOCUMENT 2 OF 3

SERVICE LEVEL AGREEMENT

Availability, support, recovery objectives and service credits for cloud services

between

[PROVIDER LEGAL NAME], [registered office / company number]
(the Provider)

and

[CUSTOMER LEGAL NAME], [registered office / company number]
(the Customer)

Agreement reference: [REF] | Effective Date: [DATE] | Version: [1.0]

Forms part of the Cloud Services Master Services Agreement (Document 1). Capitalised terms not defined here have the meaning in the MSA. Percentages, times and credit levels are market-typical starting points and must be set according to the Service Model, architecture and price. Square brackets mark items to be completed.

Contents

Part A. Overview 3

Part B. Service Level Agreement 4

1. Definitions . . . 4
2. Scope, Tiers and Interpretation . . . 4
3. Measurement of Availability . . . 4
4. Exclusions . . . 5
5. Scheduled Maintenance . . . 5
6. Incident Management and Support . . . 5
7. Service Credits . . . 6
8. Backup, Recovery and Business Continuity . . . 7
9. Chronic Failure and Termination . . . 7
10. Reporting and Service Reviews . . . 7
11. Disputes and Changes . . . 7

Schedule 1. IaaS Service Levels 8

Schedule 2. PaaS Service Levels 9

Schedule 3. SaaS Service Levels 10

Schedule 4. Support Tiers 11

Part A. Overview

A.1 How this SLA works

The SLA turns the Provider's reliability promise into measurable commitments. It sets (1) how availability is measured, (2) how incidents are prioritised and how quickly the Provider must respond and fix, (3) recovery objectives for data and service restoration, (4) what the Customer receives when the Provider falls short (service credits), and (5) when repeated failure allows the Customer to terminate. Commitments differ by Service Model, as set out in Schedule 1 (IaaS), Schedule 2 (PaaS) and Schedule 3 (SaaS).

A.2 Summary of key commitments (default Standard tier)

Commitment	IaaS	PaaS	SaaS
Monthly availability target (measured per Service)	99.95% (single-zone instance: 99.5%; multi-zone: 99.99%)	99.9%	99.9%
Storage durability (object storage)	99.999999999% annual (11 nines)	n/a	n/a
Severity 1 response / target restoration	15 min / 4 hours	30 min / 4 hours	30 min / 4 hours
Recovery Point Objective (RPO)	Customer-configured; managed backups: 24 h	15 min (managed databases)	1 hour
Recovery Time Objective (RTO)	Per architecture; managed services: 4 h	4 hours	4 hours
Maximum service credit per month	30% of monthly Fees for the affected Service	30%	30%
Chronic Failure termination right	Yes (Clause 9)	Yes	Yes

Drafting note: *Values are illustrative and should be aligned with the architecture sold (single versus multi-zone), the price tier, and the Customer's criticality. Higher tiers (Premium, Critical) typically add shorter response times and higher credits.*

Part B. Service Level Agreement

1. Definitions

"Availability" means the percentage of Measurement Minutes in a Measurement Period during which the Service is Available, calculated under Clause 3.

"Available" means that the Service is operational and accessible to the Customer, and that the Customer can perform the core functions described for that Service in the Documentation, as measured by the Provider's monitoring and confirmed by independent synthetic probes at [interval].

"Downtime" means a Measurement Minute in which the Service is not Available, excluding Excluded Events.

"Excluded Events" means the events listed in Clause 4.

"Incident" means any event that is not part of the standard operation of a Service and that causes, or may cause, an interruption or reduction in quality of that Service.

"Measurement Period" means each calendar month.

"Measurement Minute" means each minute in a Measurement Period, minus minutes of Scheduled Maintenance within the allowance in Clause 5.

"Monthly Fees" means the Fees paid or payable for the affected Service for the calendar month in which the failure occurs (or one-twelfth of annual Fees).

"RPO" means Recovery Point Objective: the maximum period of data loss, measured in time before the Incident, that the Provider commits to.

"RTO" means Recovery Time Objective: the maximum time within which the Provider restores the Service after a disaster or failure.

"Scheduled Maintenance" means planned maintenance notified under Clause 5.

"Service Credit" means the amount credited to the Customer under Clause 7.

"Severity Level" means the priority assigned to an Incident under Clause 6.

2. Scope, Tiers and Interpretation

2.1 This SLA applies to each Service in the Order Form and to the support tier stated there ([Standard / Premium / Critical]). Where the Order Form lists several Services, the Availability of each is measured separately.

2.2 The Service Level Schedules (1 to 3) apply according to the Service Model. Where a Service is composed of multiple components, the Availability commitment applies to the end-to-end Service unless the Schedule states otherwise.

2.3 This SLA prevails over the MSA in relation to service levels, support and credits only; in all other respects the MSA applies (MSA Clause 2.2).

3. Measurement of Availability

3.1 Availability for a Measurement Period is calculated as: **Availability % = ((Measurement Minutes - Downtime Minutes) / Measurement Minutes) x 100**.

3.2 Worked example: a 30-day month has 43,200 minutes. With 60 minutes of Scheduled Maintenance within the allowance, Measurement Minutes are 43,140. If there are 90 minutes of Downtime, Availability is (43,140 - 90) / 43,140 x 100 = 99.79%. Against a 99.9% target, the Service has missed the target and Service Credits apply.

3.3 For reference, the maximum Downtime permitted per 30-day month at various targets is approximately: 99.5% - 3 hours 36 minutes; 99.9% - 43 minutes; 99.95% - 21.6 minutes; 99.99% - 4.3 minutes.

3.4 The Provider shall measure Availability using its monitoring systems and shall provide the Customer with a monthly report, including raw monitoring data on request. The Customer may submit its own evidence of Downtime. If the Parties' measurements differ materially, they shall investigate jointly and, failing agreement, the matter is resolved under Clause 11.

3.5 Downtime starts when the Provider detects (or should have detected, using the monitoring described in Clause 3.1) the Incident or when the Customer reports it, whichever is earlier, and ends when the Service is restored to the Available state.

4. Exclusions

4.1 The following (**Excluded Events**) are not counted as Downtime, to the extent they actually cause the unavailability:

(a) Scheduled Maintenance within Clause 5;

(b) the Customer's acts or omissions in breach of the MSA, including misconfiguration of resources it controls under the Shared Responsibility Matrix, exceeding documented quotas or rate limits, or use inconsistent with the Documentation;

(c) failure of the Customer's own systems, networks or third-party services that are not part of the Services or are not under the Provider's or its subcontractors' control;

(d) force majeure under MSA Clause 23 (which, for the avoidance of doubt, does not include failures of the Provider's redundancy, subcontractors, or security measures);

(e) suspension of the Services in accordance with MSA Clause 8; or

(f) beta, preview or free-tier features expressly identified as outside this SLA.

4.2 The Provider has the burden of showing that an Excluded Event applies and must provide supporting evidence.

5. Scheduled Maintenance

5.1 The Provider shall perform Scheduled Maintenance only in the maintenance windows of [Saturday 22:00 to Sunday 04:00 Central European Time], give at least [seven (7)] days' notice, aim to avoid Downtime through redundancy and rolling updates, and not exceed [4] hours of maintenance-related Downtime per month and [12] hours per year.

5.2 Emergency maintenance required to address a critical security vulnerability may be performed on shorter notice, with as much notice as practicable, and counts as Downtime unless it is necessary to protect against a vulnerability not caused by the Provider's failure to patch within the periods in Schedule 4 of the MSA.

6. Incident Management and Support

6.1 **Support channels and hours.** Support is available [24x7 for Severity 1 and 2 Incidents and 08:00-18:00 CET on working days for Severity 3 and 4] via [portal, telephone, email]. The Customer shall designate [number] authorised support contacts.

6.2 **Severity classification.** The Customer proposes the Severity Level reasonably and the Provider may reclassify on reasoned grounds in consultation with the Customer.

Severity	Definition	Response (Standard)	Update frequency	Target restoration (workaround or fix)
1 - Critical	Service unavailable or unusable; data loss or corruption; confirmed Security Incident; no workaround	[30] min (24x7)	Every [30] min	[4] hours
2 - High	Major function seriously degraded; significant business impact; workaround complex	[1] hour (24x7)	Every [2] hours	[8] hours

Severity	Definition	Response (Standard)	Update frequency	Target restoration (workaround or fix)
3 - Medium	Minor function impaired; workaround available	[4] business hours	Daily	[3] business days
4 - Low	General queries, cosmetic issues, requests for information	[1] business day	As agreed	Next release / as agreed

6.3 **Response** means that a suitably qualified engineer has acknowledged the Incident, begun diagnosis and contacted the Customer. **Target restoration** means that the Service has been restored or an acceptable workaround provided, with a permanent fix to follow within [30] days.

6.4 **Escalation.** If the response or restoration target is missed, or at the Customer's request, the Incident escalates according to Schedule 6 of the MSA. Severity 1 Incidents are escalated to the Provider's duty manager automatically after [30] minutes without progress.

6.5 **Root cause analysis.** For each Severity 1 Incident and each Incident causing a failure to meet a target, the Provider shall give a written root-cause analysis within [five (5)] working days and a remediation plan with dates.

6.6 **Status page and communications.** The Provider shall maintain a public or customer-accessible status page showing real-time Service health and shall proactively notify the Customer's named contacts of Severity 1 and 2 Incidents within [15] minutes of detection.

7. Service Credits

7.1 If Availability in a Measurement Period falls below the target for a Service, the Customer is entitled to a Service Credit, as a percentage of Monthly Fees for the affected Service, as follows (for a 99.9% target):

Monthly Availability	Service Credit (% of Monthly Fees)
99.9% or above	None
Below 99.9% but at least 99.0%	10%
Below 99.0% but at least 95.0%	25%
Below 95.0%	50% (capped at the monthly maximum in Part A)

> ***Drafting note:*** *Credit schedules should be adapted for different targets. A common approach is tiers relative to the target (for example: below target by up to 0.5 points - 10%; up to 1.5 points - 25%; more than that - 50%). Customers should push for credits at a meaningful level (not nominal), automatic application, and for credits not to be the exclusive remedy.*

7.2 **Support credits.** For each Severity 1 Incident where the response target is missed by more than [100]%, the Customer receives a credit of [2]% of Monthly Fees for the affected Service, and where the restoration target is missed by more than [100]%, [5]% (subject to the cap).

7.3 **Recovery objectives credits.** If the Provider fails to meet an RTO or RPO in an actual disaster recovery event, the Customer receives a credit of [10]% of Monthly Fees for the affected Service for that month, without prejudice to other remedies.

7.4 **Claiming and application.** The Provider shall calculate and report Service Credits in its monthly report and apply them automatically to the next invoice (or pay them within [thirty (30)] days if there is no further invoice). The Customer may also claim a credit by notice within [sixty (60)] days of receiving the monthly report. Late claims are not barred where the Provider's report was missing or incomplete.

7.5 **Nature of credits.** Service Credits are a price adjustment reflecting the reduced value of the Services and not a penalty. They are [not] the sole and exclusive remedy: the Customer retains its rights to terminate for Chronic Failure, to claim damages for losses exceeding the Service Credits (but credited amounts are set off against the damages awarded for the same failure), and its other rights under the MSA, in particular for loss of Customer Data, a Security Incident, or wilful or gross-negligent breach.

7.6 Aggregate Service Credits in a month shall not exceed [thirty percent (30%)] of the Monthly Fees for the affected Service.

8. Backup, Recovery and Business Continuity

8.1 The Provider shall maintain the RPO and RTO stated in Part A and the Schedule for the Service Model, by [replication across availability zones; automated backups at the interval stated; backup retention for [30] days; encryption of backups; storage of backups in the Service Region].

8.2 The Provider shall test its backup restoration and disaster recovery plans at least annually and provide the Customer with a summary of results and remediation of failures. The Customer may participate in or witness a test of arrangements relevant to its Services, up to once a year.

8.3 For IaaS and PaaS, the Provider offers the backup and replication tools, but the Customer is responsible for configuring backups of data and configurations it controls unless a managed backup Service is ordered (MSA Schedule 3).

8.4 In a declared disaster the Provider shall prioritise restoration according to the Customer's critical-service ranking in the Order Form, and keep the Customer informed according to Clause 6.6.

9. Chronic Failure and Termination

9.1 **Chronic Failure** occurs if, for any Service: (a) Availability is below the target in any [three (3)] months within a rolling [six (6)] months; (b) Availability falls below [95%] in any one month; (c) the Provider misses the Severity 1 restoration target in [three (3)] Incidents within a rolling [six (6)] months; or (d) the Provider fails to deliver a root-cause analysis or remediation plan after two reminders.

9.2 On Chronic Failure the Customer may, by notice given within [ninety (90)] days after the last event, terminate the affected Service (or the Order Form, if the failure materially affects the whole) without early termination charges, with a refund of prepaid Fees for the period after termination, and with the support described in MSA Clause 22 (Exit). Service Credits accrued remain payable.

9.3 The Customer's termination rights under this Clause are in addition to those in MSA Clause 21.

10. Reporting and Service Reviews

10.1 The Provider shall deliver, within [ten (10)] working days after each month-end, a report containing: Availability for each Service and region; Downtime events with causes; Incident counts by Severity Level with response and restoration times; Service Credits earned; planned maintenance; capacity and performance trends; security-event summary; and open risks.

10.2 The Parties shall hold service reviews at the frequency in MSA Schedule 6 to review performance, trends and improvement actions. The Provider shall maintain a service-improvement plan for any month with a missed target.

10.3 The Customer's rights to audit or inspect the Provider's monitoring systems and records supporting these reports are as set out in MSA Clause 16.

11. Disputes and Changes

11.1 Disputes about measurement, Excluded Events or Service Credits are escalated under MSA Clause 25.2; the Customer may refer a technical dispute to an independent expert appointed by agreement (or by [institution]), whose decision on measurement is binding save for manifest error. Undisputed credits are applied in the meantime.

11.2 Service levels may not be reduced during the Term. Improvements and amendments follow MSA Clause 9. Any tier upgrade takes effect from the next Measurement Period.

Schedule 1. IaaS Service Levels

Service component	Metric	Target
Compute instance (multi-zone deployment)	Monthly Availability of at least two instances in separate zones, with the ability to run at least one	99.99%
Compute instance (single instance, premium storage)	Monthly Availability of the instance	99.9%
Compute instance (single instance, standard storage)	Monthly Availability	99.5%
Block storage / persistent volumes	Monthly Availability; durability of replicated volumes	99.9%; annual durability [99.999%]
Object storage	Monthly Availability; annual durability	99.95%; 99.999999999% (11 nines)
Virtual network / load balancer	Monthly Availability of network connectivity	99.99% (load balancer 99.99%)
Provisioning	Time to provision standard resources after valid API request	[95]% within [10] minutes
Management API / console	Monthly Availability	99.9%
Recovery	RPO / RTO for provider-managed backups and snapshots	24 h / 4 h (configurable)

IaaS-specific notes. Availability measures whether the Provider's infrastructure resource is reachable and operating, not the Customer's guest operating system or application. Where the Customer has not architected for redundancy (for example, a single instance in one zone), the single-instance target applies. Customer-induced capacity exhaustion, quota limits and failure to apply provider-notified host maintenance actions are excluded.

Schedule 2. PaaS Service Levels

Service component	Metric	Target
Application runtime / container platform	Monthly Availability of the platform control plane and runtime	99.95%
Managed database (multi-zone)	Monthly Availability	99.99% (single-zone: 99.9%)
Platform API and deployment services	Monthly Availability; successful deployment rate (platform-caused failures)	99.9%; [99]%
Messaging / event services	Monthly Availability	99.9%
Patching of managed runtimes	Critical vulnerabilities patched	Within [7] days (high: [30])
Backup / point-in-time recovery	RPO / RTO	15 min / 4 h
Deprecation notice	Minimum notice before retiring a runtime, API or version	[12] months

Schedule 3. SaaS Service Levels

Service component	Metric	Target
Application (web and API)	Monthly Availability of the production application	99.9%
Performance	[95]% of requests respond within [2] seconds at normal load	As stated
Support response	Per Clause 6.2 for each Severity Level	See table
Backup and recovery	RPO / RTO	1 h / 4 h
Release management	Notice of material changes; no production downtime for routine releases beyond the maintenance allowance	[30] days / Clause 5
Data export	Customer can export all Customer Data via interface or API at any time; bulk export delivered within	[5] working days on request
Security incidents	Notification	Within [24 / 48] hours of awareness (MSA Clause 12.2)

Schedule 4. Support Tiers

Feature	Standard	Premium	Critical
Hours	Business hours (Sev 3-4); 24x7 for Sev 1-2	24x7 all severities	24x7 all severities, dedicated engineer
Sev 1 response	30 min	15 min	10 min
Sev 1 target restoration	4 hours	2 hours	1 hour
Named service manager	Shared pool	Named	Named plus technical account manager
Service reviews	Annual	Quarterly	Monthly
Availability target uplift	None	None	[+0.05]% where the architecture supports it
Maximum monthly credit	30%	40%	50%

Signed for and on behalf of [PROVIDER LEGAL NAME]	Signed for and on behalf of [CUSTOMER LEGAL NAME]
Signature: _____ Name: [NAME] Title: [TITLE] Date: [DATE]	Signature: _____ Name: [NAME] Title: [TITLE] Date: [DATE]